

JD DAVENPORT

me@jddavenport.com | 801-694-7791 | Provo, UT | [linkedin.com/in/jd-davenport](https://www.linkedin.com/in/jd-davenport) | github.com/jddavenportOpen | jddavenport.com

Product manager who ships AI products and gets them adopted at enterprise scale. Grew National Grid's digital services revenue from \$2M to \$20M+ annualized and launched a GPT-powered RAG platform to 3,000+ sustained monthly active users. At Siemens, built Spark AI, a 120+ agent platform, and took active adoption from 3% to 65% in 10 weeks. Has built production Claude systems since November 2025, including a multi-agent OS on the Anthropic API used across 180+ production repositories.

EXPERIENCE

AI Product Strategist, XPS Program | Brightly Software / Siemens Summer 2026

- Built and launched Spark AI - custom AI platform for Siemens US comprising 120+ custom agents and 50+ custom skills, deployed across 1,000+ employees spanning all major job functions and seniority levels
- Drove 65% active adoption from a 3% baseline in 10 weeks - 650+ employees actively using by close; enterprise AI rollouts at comparable industrial companies typically take quarters and rarely exceed 30%
- Designed and presented 3/6/12-month AI transformation roadmap to Siemens US leadership - sequenced by function-specific readiness, dependency-mapped, with role-specialized success metrics tied to measurable workflow efficiency outcomes

Founder & Managing Director | Product Builder / AI Foundry at BYU 2026 - Present

- Founded AI product-building program inside BYU Marriott School of Business - identified market gap, built business case, secured faculty sponsorship and institutional backing, designed curriculum and cohort selection model from scratch, recruited and launched 30-builder founding cohort; accepted into Anthropic's Claude Partner Network
- Cohort builders are shipping real AI products with external users; curriculum integrates product discovery, agentic system design, and launch - not AI theory
- Operating at startup speed inside an academic institution whose incentives resist it; every resource requires cross-functional justification and institutional buy-in without a budget or a mandate

Principal Product Manager, Digital Services | National Grid 2022 - 2025

Fortune 500 utility, 23,000 employees, 8 fragmented business units

- Designed and deployed Universal Prioritization Framework - replacing emotion- and seniority-based build decisions with measurable criteria (user impact, business impact, technical cost/risk, strategic dependency) adopted across dozens of teams; became the shared decision language enabling cross-BU collaboration for the first time
- Founded Cloud Leadership Sync - first executive and Chief Architect-level cross-BU coordination forum in company history; introduced portfolio-level roadmapping, PRDs, and agile ceremonies at scale across all business units, structurally transforming a siloed engineering organization into a product organization
- Launched GPT-powered RAG platform achieving 3,000+ sustained monthly active users - delayed launch 6 weeks for document-level access controls (different BUs held confidential assets incompatible with universal-search); tracked cohort retention curves, identified drop-off cohorts, iterated on discovery; 3,000+ MAU reflects ongoing monthly activity
- Built and deployed Infrastructure as Code (IaC) service reducing cloud service deployment from weeks to minutes; standardized 100+ enterprise tools (AKS, Snowflake, Power BI) across all BUs with no separate procurement or security review per team
- Ran individual stakeholder campaigns convincing every BU President and Chief Architect to cede infrastructure autonomy - pitched Presidents on engineering tax elimination, Chief Architects on production-grade tooling with zero security review burden
- Converted skeptical data science team from "product management is overhead" to "product management is why our models get used" - through embedded discovery and their own data, not authority
- Scaled engineering organization from 5 to 35 engineers; drove digital services revenue from \$2M to \$20M+ annualized
- Promoted: Technical PM -> Senior PM (18 months) -> Principal PM (6 months) - top-performer track

Senior Product Consultant | Deloitte

2020 - 2022

- Led Pfizer's contact-center AI engagement, setting what the model answered and what went to a person - average queue time fell from 75 minutes to 15
- Found \$17.5M of configuration savings at GM and presented it as the analytics lead
- Led internal platform and technology transformation engagements for Pfizer, GM, and Walmart at the VP and C-suite level
- Promoted to Senior Consultant in year one - top 5% of hiring cohort

Senior Full-Stack Developer & Product Owner | Brio Energy

2018 - 2020

- Built custom CRM from scratch; drove \$3.9M in new revenue through marketing platform rebuild

Full-Stack Developer | Surge Software

2017 - 2018

- Built 27 landing pages generating 17M visits; shipped 12 APIs in 3.5 months

AI PRODUCTION SYSTEMS

Built Nerve Center - a production implementation of org-scale Claude orchestration running since November 2025: Claude Opus orchestrator routing to 10+ domain-isolated agents (each scoped by function and data boundary, with differentiated access and permission models), a durable inter-agent messaging bus (Supabase-backed, dead-letter/retry/alerts), self-healing infrastructure with integrity guards every 15 minutes, and a real-time web cockpit with audit trail across all agents. Backed by Anthropic API across 180+ production repositories - production reliability requirements, not experimentation: 10 automated domains create real failure-point exposure, not personal convenience. Architecture designed to adopt new model capabilities without rebuilding the deployment layer.

Operational knowledge of Claude's failure modes that documentation doesn't cover: context window degradation curves, multi-agent state divergence patterns, prompting architecture differences across agent-loop vs. single-turn vs. long-running contexts. Working implementation of AI evaluation loops: cohort-level performance tracking, failure classification by agent type, and routing adjustments driven by observed degradation patterns.

Has designed the architectural shift required for org-scale Claude deployment: single-user full-trust model breaks immediately at organizational scale, which requires differentiated access and permission models, safety and privacy controls, data handling tradeoffs, audit trails, and opt-out capability; shared organizational context (Claude has no default knowledge of users or their work) requires active discovery and synthesis architecture; and reliability requirements shift categorically when an outage blocks a whole team rather than one person. Each is already solved for a single user; the product work is making it hold for an entire organization.

EDUCATION

BYU Marriott School of Business

2025-2027

MBA, Strategy & Product Management

One class from completion

Georgia Tech

2020 - 2022

Graduate Coursework, Data Science (OMSA)

Brigham Young University

B.S., Management Information Systems

TECHNICAL SKILLS

AI Systems: Claude API, multi-agent orchestration, RAG system design, enterprise AI deployment, agentic workflow design, prompt architecture for production agents, SSO/identity integration, AI evaluation loops

Infrastructure: Azure AKS, Snowflake, Power BI, Supabase, PostgreSQL, Google Workspace (Docs, Drive, Gmail, Calendar), Infrastructure as Code

Development: Python, TypeScript, full-stack development, REST APIs

Product: portfolio roadmapping, PRDs, OKRs, agile at scale, stakeholder alignment without authority, cross-functional discovery, org-scale AI deployment, AI product strategy, enterprise AI adoption